

Hosting Qwen2.5/Qwen3 on a VPS using vLLM

1. VPS Requirements

Ubuntu 22.04, NVIDIA GPU, CUDA drivers, Python 3.10+, pip, and sufficient VRAM (16GB+ recommended for 3B/7B models).

2. Install Dependencies

```
sudo apt update
pip install torch vllm transformers huggingface_hub
```

3. Download Model

Download the fine-tuned Qwen model from Hugging Face or copy the trained adapter/model to the VPS.

4. Start vLLM

Run:

```
python -m vllm.entrypoints.openai.api_server --model ./qwen_astrologer --host 0.0.0.0 --port 8000
```

5. Test API

Send an OpenAI-compatible POST request to `/v1/chat/completions` with system and user messages.

6. Production Deployment

Use Nginx as a reverse proxy, enable HTTPS using Let's Encrypt, run vLLM as a systemd service, enable firewall rules, and monitor logs.